

# EMBODIED MODEL REPAIR MEMORY: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

**Anonymous authors**

Paper under double-blind review

## ABSTRACT

We study adaptive world models for long-horizon robot tasks. The motivating bet is: Store repairs to model beliefs as reusable state, not transient corrective prompts.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

## 1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For adaptive world models for long-horizon robot tasks, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

## 2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026; 2019; 2024; 2025a; 2021; 2025b). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

## 3 MECHANISM

Given state  $s$  and action  $a$ , the model returns an observed next-state prediction plus a finite atlas  $B(s, a)$  of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

## 4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

| Variant                   | Split             | Success | Tail risk |
|---------------------------|-------------------|---------|-----------|
| observed_only_wam         | nominal           | 0.590   | 0.410     |
| observed_only_wam         | hidden_mode_shift | 0.165   | 0.835     |
| augmentation_as_data_wam  | nominal           | 0.608   | 0.392     |
| augmentation_as_data_wam  | hidden_mode_shift | 0.233   | 0.767     |
| scalar_uncertainty_wam    | nominal           | 0.685   | 0.315     |
| scalar_uncertainty_wam    | hidden_mode_shift | 0.305   | 0.695     |
| proposed_branch_mechanism | nominal           | 0.762   | 0.238     |
| proposed_branch_mechanism | hidden_mode_shift | 0.605   | 0.395     |

## 5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

## 6 CONCLUSION

Embodied Model Repair Memory argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

## REFERENCES

- Scene memory transformer for embodied agents in long-horizon tasks, 2019. <https://doi.org/10.1109/cvpr.2019.00063>.
- Calvin: A benchmark for language-conditioned policy learning for long-horizon robot manipulation tasks, 2021. <http://arxiv.org/abs/2112.03227v4>.
- From real world to logic and back: Learning generalizable relational concepts for long horizon robot planning, 2024. <http://arxiv.org/abs/2402.11871v6>.
- Map-vla: Memory-augmented prompting for vision-language-action model in robotic manipulation, 2025a. <http://arxiv.org/abs/2511.09516v1>.
- Furniturebench: Reproducible real-world benchmark for long-horizon complex manipulation, 2025b. <https://doi.org/10.1177/02783649241304789>.
- Pi-vla: Adaptive symmetry-aware decision-making for long-horizon vision-language-action manipulation, 2026. <https://doi.org/10.3390/sym18030394>.